

EMB-Lab, Fak. N

Convolutional Neural Networks

DLM – Deep Learning Methoden


Piotr Bialecki

Carla Gil, Benjamin Kraus, Prof. Dr. Marcus Vetter


Mannheim, 27.11.2017




- (Deep) Neural Networks
 - Layer connectivity
 - Nonlinearities
 - Vanishing / Exploding Gradient
- Hyperparamters
 - Learning rate
 - ...



(a) ReLU




(b) LeakyReLU



[Karpathy, CS231n]




- Loop:
 - **Sample** a batch of data
 - **Forward** prop it through the network, get loss
 - **Backprop** to calculate the gradients
 - **Update** the parameters (weights) using the gradient



[Karpathy, CS231n]

Can we use Fully-connected nets for image classification?

Is a simple reshaping of the image working?




[Karpathy, CS231n]



Images are (often) large

- MNIST: $28 \times 28 \times 1$ (width x height x channels)
= 784 input dimensions




- $200 \times 200 \times 3$
= 120,000 input dimensions



leopard


- $512 \times 512 \times 3$
= 786,432 input dimensions






Feeding the reshaped image into a fully connected net yields a huge amount of parameters (weights)

- Weight matrix: $W = w_{ij} \in \mathbb{R}^{m \times n}$
- With m pixels (input) and n neurons in the 1st hidden layer




[Yoshi Komiri]

- Huge amount of weights leads to large **network capacity**
- Each image can be understood as a single point in a high-dimensional feature space
- Remember the **curse of dimensionality**
 - The volume of the feature space increases so fast that the available data become sparse
 - Amount of data to "fill" the space grows exponentially



[Bishop, Pattern Recognition and Machine Learning]

- Human intuition is mostly wrong in high dimensions!
- E.g.: Where do you think is the most probability mass of a high dimensional multivariate normal distribution?






- Human intuition is mostly wrong in high dimensions!
- E.g.: Where do you think is the most probability mass of a high dimensional multivariate normal distribution?
- -> It's concentrated in a thin shell some distance away from the origin!






- Human intuition is mostly wrong in high dimensions!
- E.g.: Where do you think is the most probability mass of a high dimensional multivariate normal distribution?
- -> It's concentrated in a thin shell some distance away from the origin!
- Volume of a sphere in d dimensions is proportional to r^d .
 - For delta radius: $\text{vol_shell} = \text{delta_r} * d * r^{(d-1)}$
- Probability density: $\text{prob_dens} = \exp(-r^2 / 2)$
- Therefore: Probability mass: $\text{vol_shell} * \text{prob_dens}$
- --> Leads to Chi distribution






Images have (often) a local structure

- Not every pixel contains new "information"
 - Neighboring pixels are correlated
 - Using every pixel is wasteful
- Image data can be seen as a composition of semantic elements
 - e.g. we only need to learn the leopard pattern once



The nature as a role model



- Hubel & Wiesel
 - 1959: Receptive Fields of Single Neurones in the Cat's Striate Cortex
 - 1962: Receptive Fields, Binocular Interaction and Functional Architecture in the Cat's Visual Cortex
 - 1968: ...






The nature as a role model

- Hubel & Wiesel's main findings are that the visual cortex is organized hierarchically
- Low level features (edges, ...)
- To high level "concepts"



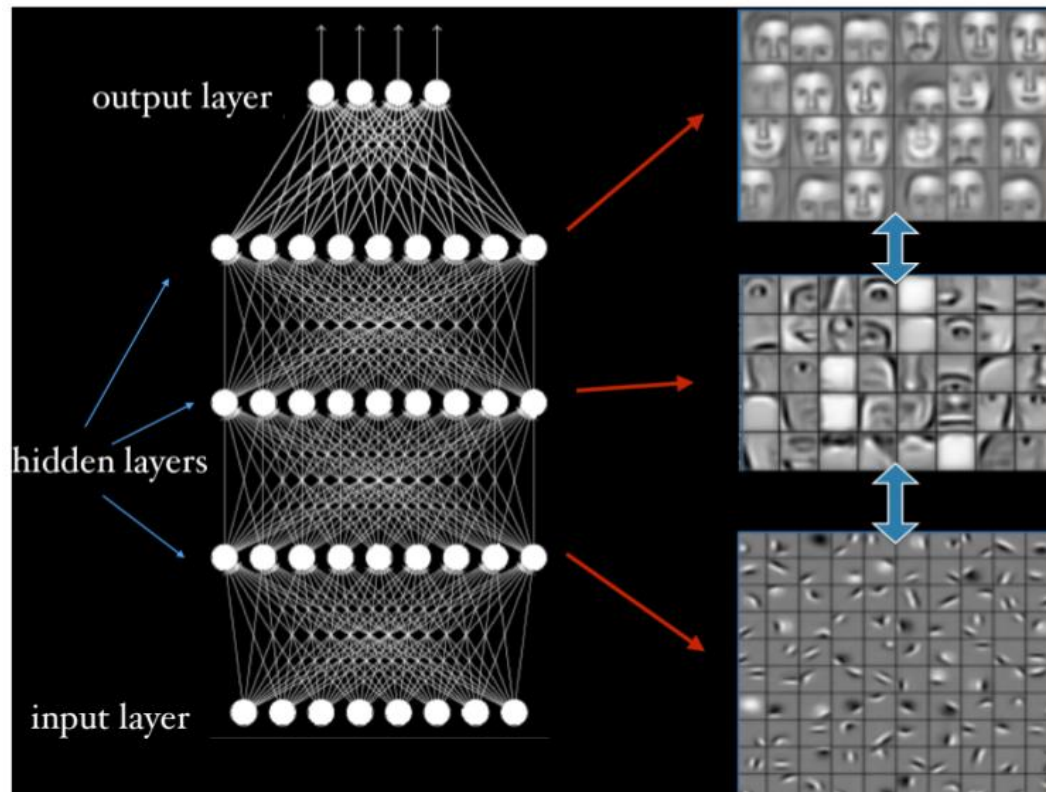


[Yann LeCun 1980, LeNet-5]



ConvNets imitate the "mammalian" vision

- Early layers learn "simple" features
- Later layers stack these features together and produce more complex features




[Lee et al.,
Convolutional deep
belief networks for
scalable
unsupervised
learning of
hierarchical
representations]



ConvNets imitate the "mammalian" vision

- And perform very well



[NVIDIA]




ConvNets consist of

- *Convolution Layers*
- *Subsampling/ Pooling Layers*
- Detector Layer: Rectified Linear Units (Nonlinearities)
- Fully-Connected Layers



Convolution layers


- Are the core building blocks of ConvNets
- Consist of a set of learnable filters
- Each filter is small spatially (width and height) but extends through the full depth of the input volume






Convolution layers

- Filters slide through the input activation and compute a dot product of its filter weights and the input at any position
- Same as convolution of input and filter



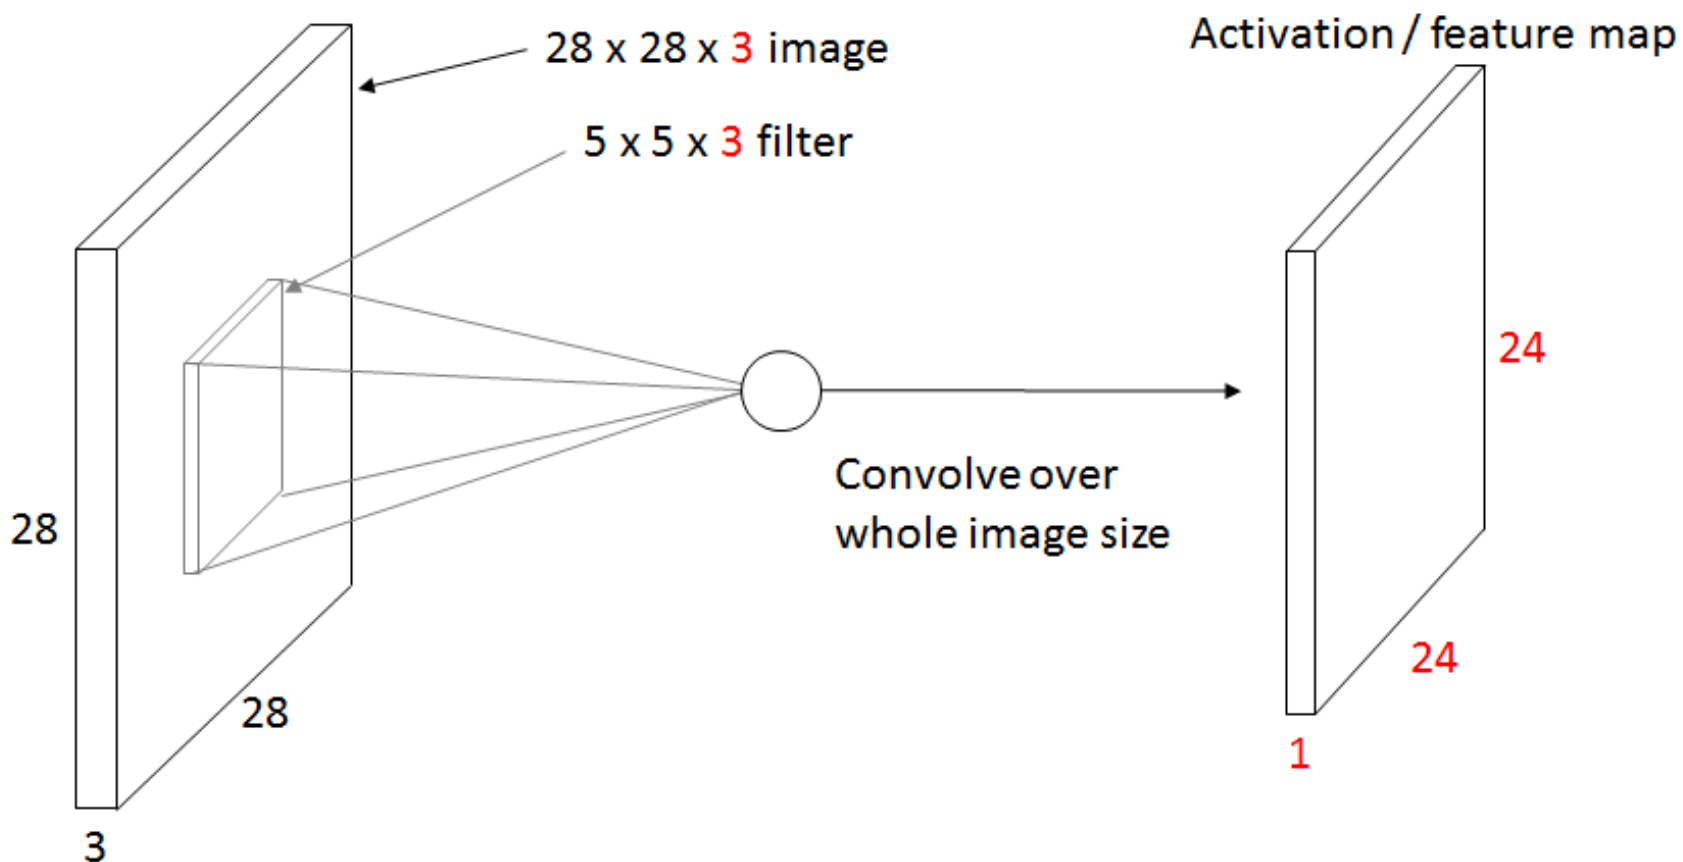


Convolution layers





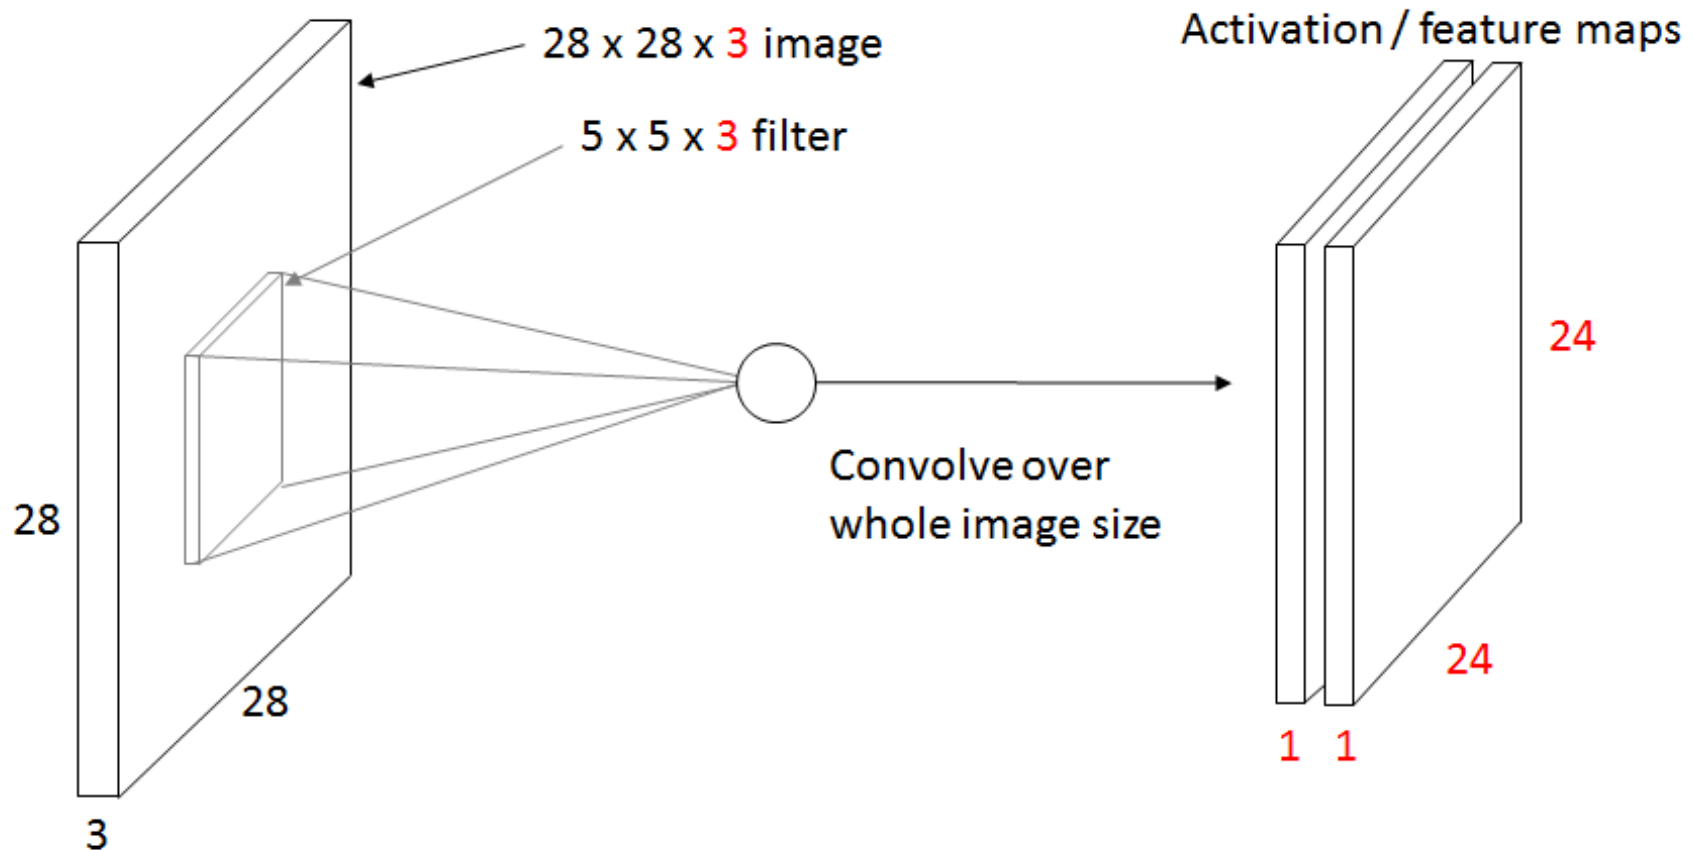
Convolution layers





Convolution layers

A second filter creates another feature map!

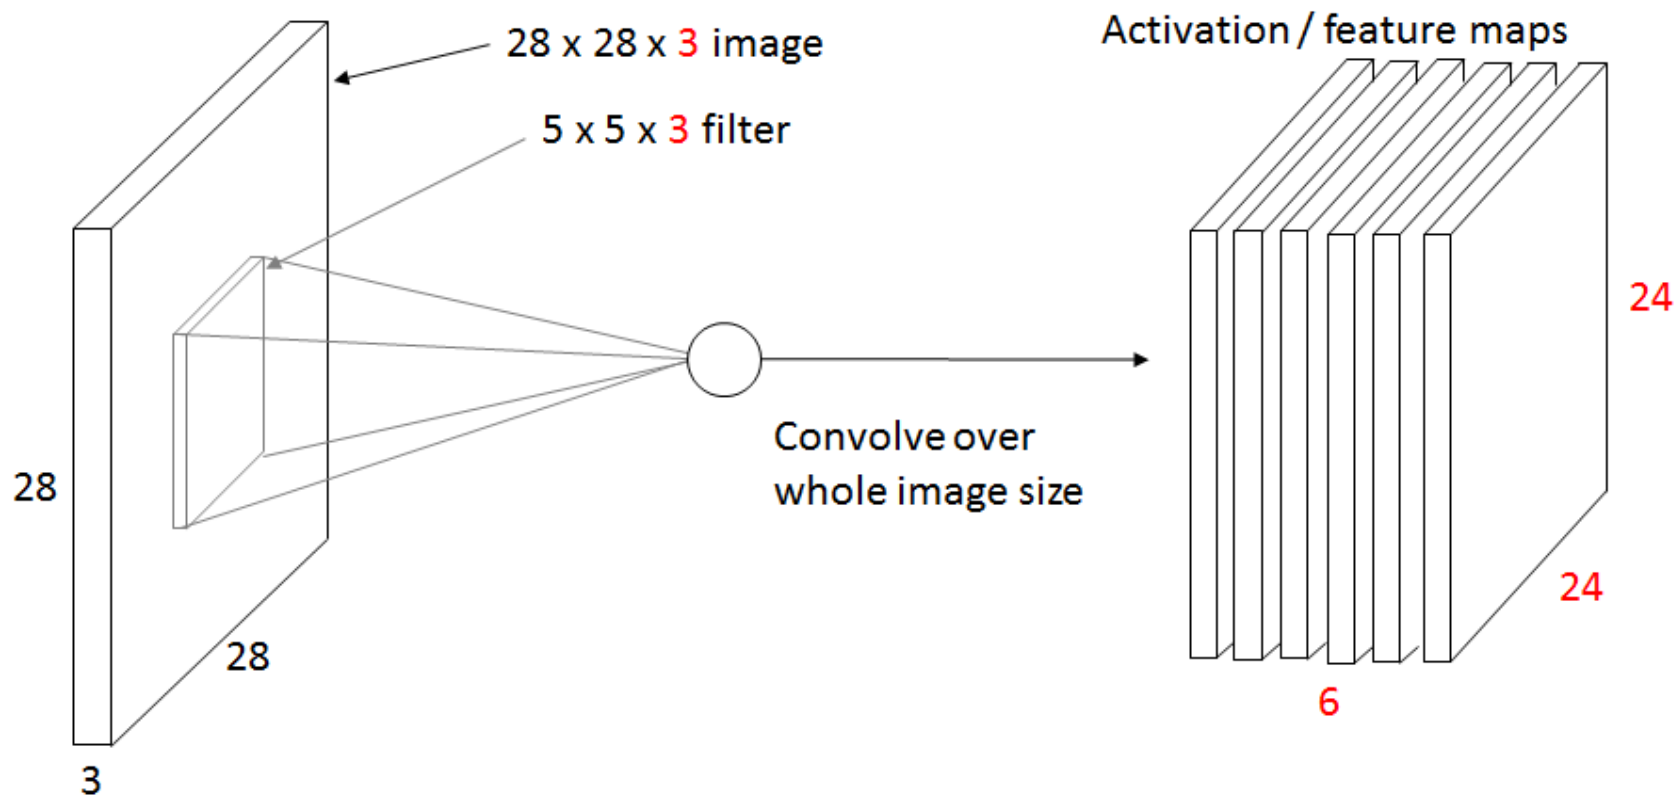




Convolution layers


6 5×5 filters create 6 separate feature maps.

These feature maps are stacked into a $24 \times 24 \times 6$ feature volume / image





A ConvNet is a sequence of ConvLayers connected with activation functions



[Karpathy, CS231n]



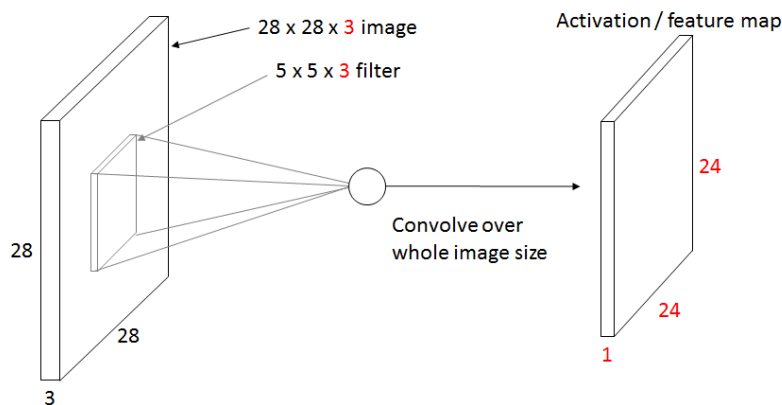
Convolutions affect spatial dimensions

- General formula $s(t) = \int x(a)w(t-a)da$

[Goodfellow et al.,
Deep Learning]

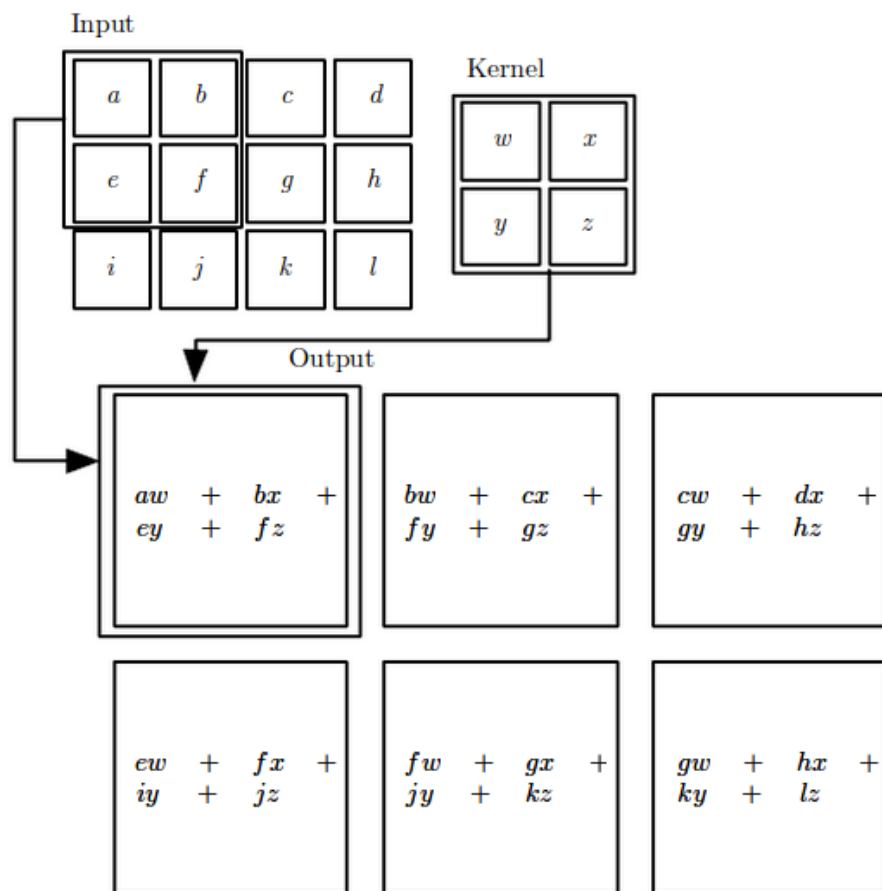
- Often denoted as $s(t) = (x * w)(t)$

- Discrete formula $s(t) = (x * w)(t) = \sum_{a=-\infty}^{\infty} x(a)w(t-a)$





"Valid" 2-D convolution without kernel flipping




[Goodfellow et al.,
Deep Learning]



"Valid" 2-D convolution without kernel flipping

7



7


7x7 input (spatially)
assume 3x3 filter

[Karpathy, CS231n]



"Valid" 2-D convolution without kernel flipping

7



7


7x7 input (spatially)
assume 3x3 filter

[Karpathy, CS231n]



"Valid" 2-D convolution without kernel flipping

7



7


7x7 input (spatially)
assume 3x3 filter

[Karpathy, CS231n]



"Valid" 2-D convolution without kernel flipping

7



7


7x7 input (spatially)
assume 3x3 filter

[Karpathy, CS231n]



"Valid" 2-D convolution without kernel flipping

7



7x7 input (spatially)
assume 3x3 filter


=> 5x5 output

7

[Karpathy, CS231n]



"Valid" 2-D convolution without kernel flipping, stride 2




7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**

[Karpathy, CS231n]



"Valid" 2-D convolution without kernel flipping, stride 2




7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**

[Karpathy, CS231n]



"Valid" 2-D convolution without kernel flipping, stride 2




7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**
=> 3x3 output!

[Karpthy, CS231n]



"Valid" 2-D convolution without kernel flipping



Output size:
 $(N - F) / \text{stride} + 1$

e.g. $N = 7, F = 3$:

stride 1 $\Rightarrow (7 - 3) / 1 + 1 = 5$

stride 2 $\Rightarrow (7 - 3) / 2 + 1 = 3$

stride 3 $\Rightarrow (7 - 3) / 3 + 1 = 2.33 \Rightarrow 2$

[Karpathy, CS231n]



Zero padding of border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

7x7 output!

[Karpathy, CS231n]



Zero padding of border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

7x7 output!

in general, common to see CONV layers with stride 1, filters of size $F \times F$, and zero-padding with $(F-1)/2$. (will preserve size spatially)

e.g. $F = 3 \Rightarrow$ zero pad with 1

$F = 5 \Rightarrow$ zero pad with 2

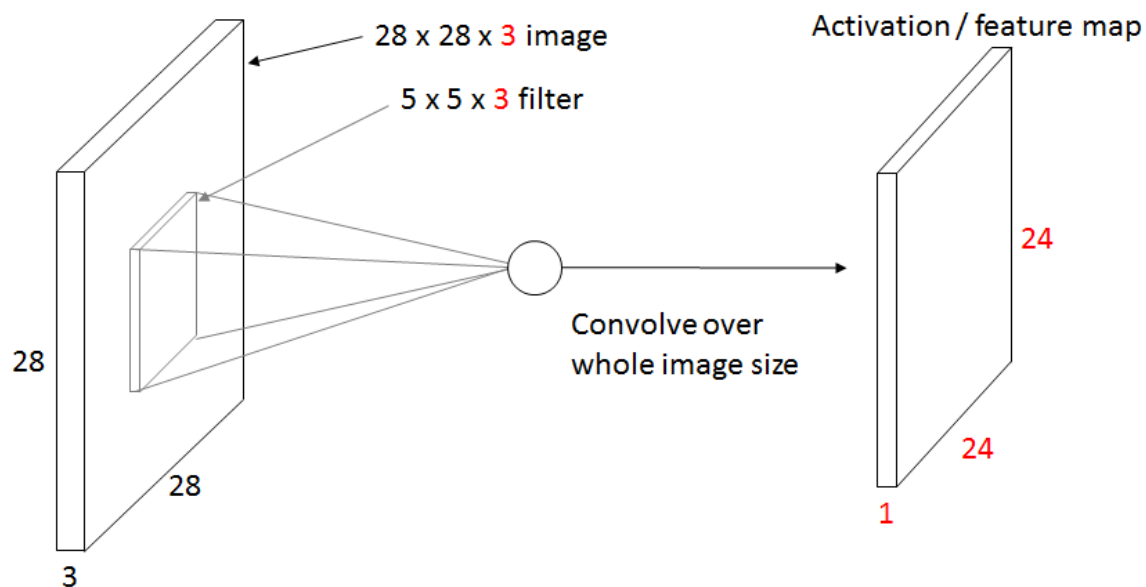
$F = 7 \Rightarrow$ zero pad with 3

[Karpathy, CS231n]



Zero padding of border might preserve the size

Otherwise the volume shrinks spatially!





Zero padding of border might preserve the size

Example:



Input image: 28x28x3

10 5x5 filters with stride **1**, pad **2**

Output volume size:

???

28 x 28 x 3 image



Output size:

$$(N - F) / \text{stride} + 1$$

e.g. $N = 7$, $F = 3$:

$$\text{stride } 1 \Rightarrow (7 - 3) / 1 + 1 = 5$$

$$\text{stride } 2 \Rightarrow (7 - 3) / 2 + 1 = 3$$

$$\text{stride } 3 \Rightarrow (7 - 3) / 3 + 1 = 2.33 \text{ :}\backslash$$



Zero padding of border might preserve the size

Example:

Input image: 28x28x3

10 5x5 filters with stride **1**, pad **2**

Output volume size:

$$(28 + 2 \cdot 2 - 5) / 1 + 1$$



$$(28 + 4 - 5) / 1 + 1$$

$$27 + 1 = 28 \text{ spatially}$$

Output volume

28x28x10

28 x 28 x 3 image



Output size:

$$(N - F) / \text{stride} + 1$$

e.g. $N = 7, F = 3$:

$$\text{stride } 1 \Rightarrow (7 - 3) / 1 + 1 = 5$$

$$\text{stride } 2 \Rightarrow (7 - 3) / 2 + 1 = 3$$

$$\text{stride } 3 \Rightarrow (7 - 3) / 3 + 1 = 2.33 \text{ :}\backslash$$



Number of parameters

Example:


Input image: $28 \times 28 \times 1$

10 5×5 filters with stride 1, pad 2

Number of parameters in layer?

???

$28 \times 28 \times 3$ image





Number of parameters

Example:

Input image: 28x28x1


10 5x5 filters with stride 1, pad 2

Number of parameters in layer?

Each filter has $5*5*1+1 = 26$ params (+1 for bias)

--> total $26 * 10 = \mathbf{260}$

28 x 28 x 3 image





Number of parameters


Example:

Input image: $28 \times 28 \times 3$

10 5×5 filters with stride 1, pad 2

Number of parameters in layer?

$28 \times 28 \times 3$ image





Number of parameters

Example:

Input image: 28x28x3


10 5x5 filters with stride 1, pad 2

Number of parameters in layer?

Each filter has $5 \times 5 \times 3 + 1 = 76$ params (+1 for bias)

--> total $76 \times 10 = \mathbf{760}$

28 x 28 x 3 image





Summary. To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a valid convolution of the d -th filter over the input volume with a stride of S , and then offset by d -th bias.

[Karpathy, CS231n]




- Keras Conv2D layer

- **filters**: Integer, the dimensionality of the output space (i.e. the number output of filters in the convolution).
- **kernel_size**: An integer or tuple/list of 2 integers, specifying the width and height of the 2D convolution window. Can be a single integer to specify the same value for all spatial dimensions.
- **strides**: An integer or tuple/list of 2 integers, specifying the strides of the convolution along the width and height. Can be a single integer to specify the same value for all spatial dimensions. Specifying any stride value $\neq 1$ is incompatible with specifying any `dilation_rate` value $\neq 1$.
- **padding**: one of `"valid"` or `"same"` (case-insensitive).
- **data_format**: A string, one of `channels_last` (default) or `channels_first`. The ordering of the dimensions in the inputs. `channels_last` corresponds to inputs with shape `(batch, height, width, channels)` while `channels_first` corresponds to inputs with shape `(batch, channels, height, width)`. It defaults to the `image_data_format` value found in your Keras config file at `~/.keras/keras.json`. If you never set it, then it will be `"channels_last"`.
- **dilation_rate**: an integer or tuple/list of 2 integers, specifying the dilation rate to use for dilated convolution. Can be a single integer to specify the same value for all spatial dimensions. Currently, specifying any `dilation_rate` value $\neq 1$ is incompatible with specifying any stride value $\neq 1$.
- **activation**: Activation function to use (see [activations](#)). If you don't specify anything, no activation is applied (ie. "linear" activation: $a(x) = x$).
- **use_bias**: Boolean, whether the layer uses a bias vector.
- **kernel_initializer**: Initializer for the `kernel` weights matrix (see [initializers](#)).
- **bias_initializer**: Initializer for the bias vector (see [initializers](#)).
- **kernel_regularizer**: Regularizer function applied to the `kernel` weights matrix (see [regularizer](#)).
- **bias_regularizer**: Regularizer function applied to the bias vector (see [regularizer](#)).
- **activity_regularizer**: Regularizer function applied to the output of the layer (its "activation"). (see [regularizer](#)).
- **kernel_constraint**: Constraint function applied to the kernel matrix (see [constraints](#)).
- **bias_constraint**: Constraint function applied to the bias vector (see [constraints](#)).



Parameter sharing

- In fully-connected net, each weight is used once to compute output
- In ConvNet each weight of the filter is used at **every** position of the input
- The weights are **tied**!




[Goodfellow et al.,
Deep Learning]



Parameter sharing

- Instead of learning a weight set for every location, we learn only one set
- Idea: A certain filter (e.g. edge detector) might be useful for every location in the input volume




[Goodfellow et al.,
Deep Learning]



Each filter learn a different feature

There will be 5 different neurons / filters
looking at the same location in the input volume




[Karpathy, CS231n]



Each filter learn a different feature

- We get 5 **depth slices** of size 28x28
- During backprop, every neuron in the volume will compute gradient for its weights
- Gradients will be added up across each depth slice and only update a single set of weights per slice





[Karpathy, CS231n]



Each filter learn a different feature

1st layer features of learned ConvNet






[Zeiler and Fergus]



Gabor filters



- Linear filters used for edge detection
- Gaussian kernel function modulated by a sinusoidal plane wave
- Simple cells in visual cortex of mammalian brains can be modeled by Gabor functions



[Goodfellow et al.,
Deep Learning]




Gabor-like learned Kernels



[Goodfellow et al.,
Deep Learning]



- The receptive field of the units in the deeper layers of a ConvNet gets larger
- This effect can be increased using strided convolutions or *pooling*
- Thus units in deeper layers can be connected to all or most of the input image




[Goodfellow et al.,
Deep Learning]



Pooling layers

- Makes the feature maps smaller and more manageable
- Operates over each feature map independently




[Karpathy, CS231n]



Pooling layers

- Replaces the ConvLayer -> ReLU output with a summary statistic of the nearby outputs

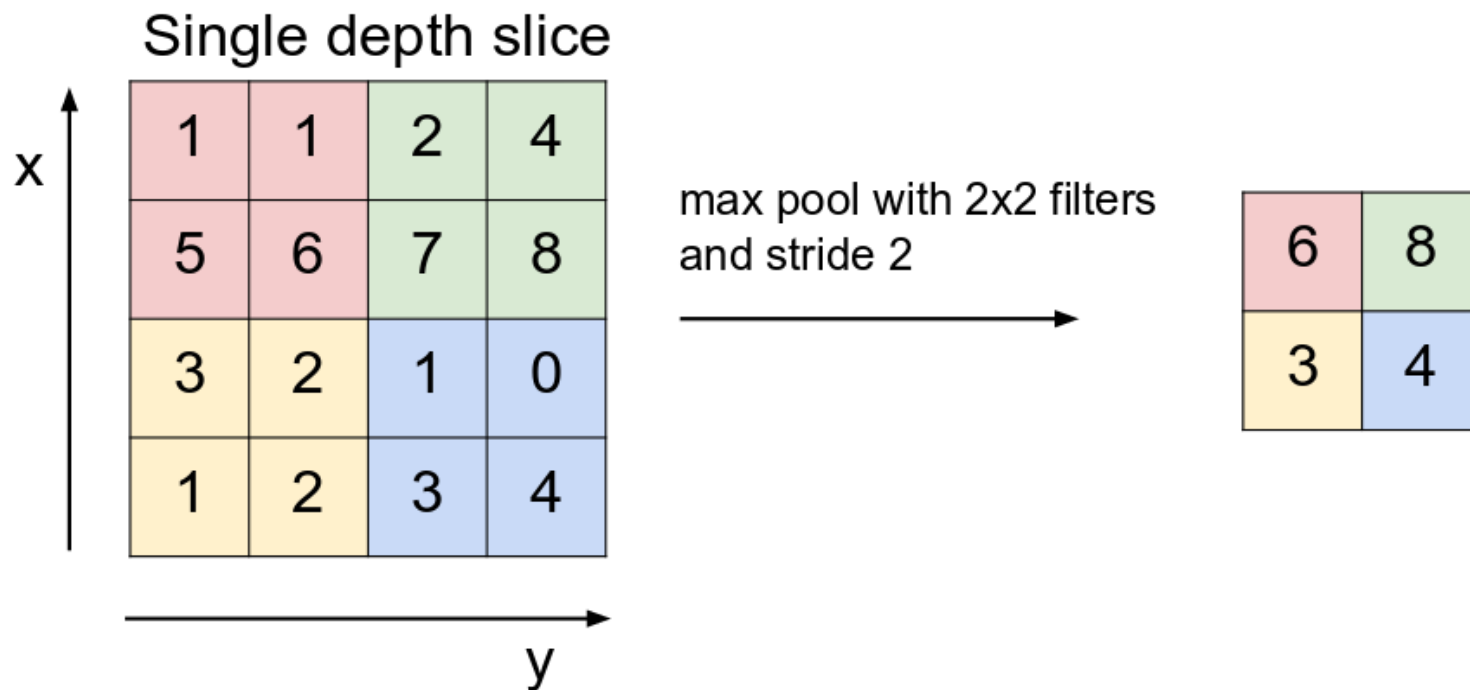


[Karpathy, CS231n]



Pooling layers

- Max Pooling




[Karpathy, CS231n]



Pooling layers

- Max Pooling
- Eliminates 75% of the feature map



[Karpathy, CS231n]



Pooling layers

- Max Pooling



- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires three hyperparameters:
 - their spatial extent F ,
 - the stride S ,
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F)/S + 1$
 - $H_2 = (H_1 - F)/S + 1$
 - $D_2 = D_1$
- Introduces zero parameters since it computes a fixed function of the input
- Note that it is not common to use zero-padding for Pooling layers

[Karpathy, CS231n]



Pooling layers

- Pooling helps to make features become approx. **invariant** to small translations of the input



[Goodfellow et al.,
Deep Learning]



Pooling layers

- Pooling helps to make features become approx. **invariant** to small translations of the input

Why is this interesting?

- Can be useful when exact location of a feature is not important
- E.g. in Face detection
- Not necessary to know pixel-perfect location of eyes, nose, mouth, ...
- Better: One eye on left side of face, another on right side

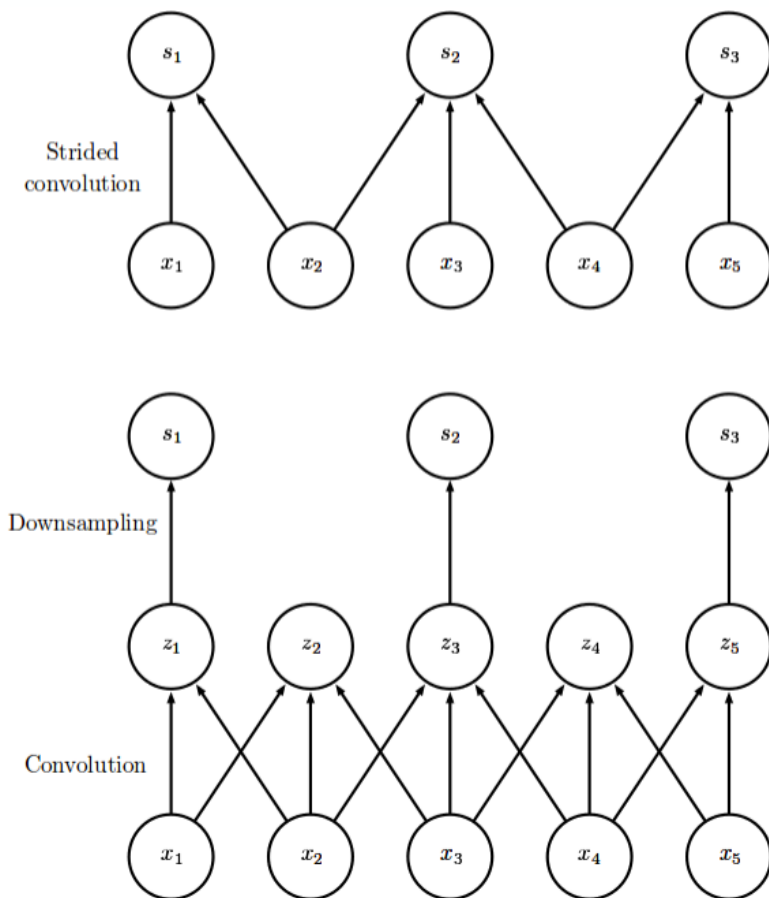


Pooling layers

- **Max Pooling**
 - Average Pooling
 - L2-norm Pooling
- --> However, pooling can also be achieved using strided convolutions!
- You could and should test both approaches!



Convolution with Stride




[Goodfellow et al.,
Deep Learning]



Common architectures


- Stacking of some Conv-ReLU layers
- Followed by a Pooling layer
- Repeat until Feature map has small size
- Add Fully-Connected layer
- Last Fully-connected layer holds outputs (scores)



[Karpathy, CS231n]

Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

First layer (CONV1): 96 11x11 filters applied at stride 4


=>

Q: what is the output volume size? Hint: $(227-11)/4+1 = 55$

[Karpathy, CS231n]

Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

First layer (CONV1): 96 11x11 filters applied at stride 4

=>


Output volume **[55x55x96]**

Q: What is the total number of parameters in this layer?

[Karpathy, CS231n]

Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

First layer (CONV1): 96 11x11 filters applied at stride 4

=>


Output volume **[55x55x96]**

Parameters: $(11*11*3)*96 = 35K$

[Karpathy, CS231n]

Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

After CONV1: 55x55x96


Second layer (POOL1): 3x3 filters applied at stride 2

Q: what is the output volume size? Hint: $(55-3)/2+1 = 27$

[Karpathy, CS231n]

Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

After CONV1: 55x55x96

Second layer (POOL1): 3x3 filters applied at stride 2


Output volume: 27x27x96

Q: what is the number of parameters in this layer?

[Karpathy, CS231n]

Case Study: AlexNet

[Krizhevsky et al. 2012]



Input: 227x227x3 images

After CONV1: 55x55x96

Second layer (POOL1): 3x3 filters applied at stride 2

Output volume: 27x27x96

Parameters: 0!

[Karpathy, CS231n]

Case Study: AlexNet

[Krizhevsky et al. 2012]

Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1


[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)



[Karpathy, CS231n]

Case Study: VGGNet

[Simonyan and Zisserman, 2014]

Only 3x3 CONV stride 1, pad 1
and 2x2 MAX POOL stride 2

best model

11.2% top 5 error in ILSVRC 2013

->

7.3% top 5 error

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: Number of parameters (in millions).

Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

[Karpathy, CS231n]



INPUT: [224x224x3] memory: 224*224*3=150K params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: 224*224*64=3.2M params: (3*3*3)*64 = 1,728

CONV3-64: [224x224x64] memory: 224*224*64=3.2M params: (3*3*64)*64 = 36,864

POOL2: [112x112x64] memory: 112*112*64=800K params: 0

CONV3-128: [112x112x128] memory: 112*112*128=1.6M params: (3*3*64)*128 = 73,728

CONV3-128: [112x112x128] memory: 112*112*128=1.6M params: (3*3*128)*128 = 147,456

POOL2: [56x56x128] memory: 56*56*128=400K params: 0

CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*128)*256 = 294,912

CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*256)*256 = 589,824

CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*256)*256 = 589,824

POOL2: [28x28x256] memory: 28*28*256=200K params: 0

CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*256)*512 = 1,179,648

CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*512)*512 = 2,359,296

CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*512)*512 = 2,359,296

POOL2: [14x14x512] memory: 14*14*512=100K params: 0

CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296

CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296

CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296

POOL2: [7x7x512] memory: 7*7*512=25K params: 0

FC: [1x1x4096] memory: 4096 params: 7*7*512*4096 = 102,760,448

FC: [1x1x4096] memory: 4096 params: 4096*4096 = 16,777,216

FC: [1x1x1000] memory: 1000 params: 4096*1000 = 4,096,000

ConvNet Configuration			
B	C	D	
13 weight layers	16 weight layers	16 weight layers	19
put (224 × 224 RGB image)			
conv3-64	conv3-64	conv3-64	cc
conv3-64	conv3-64	conv3-64	cc
maxpool			
conv3-128	conv3-128	conv3-128	co
conv3-128	conv3-128	conv3-128	co
maxpool			
conv3-256	conv3-256	conv3-256	co
conv3-256	conv3-256	conv3-256	co
	conv1-256	conv3-256	co
			co
maxpool			
conv3-512	conv3-512	conv3-512	co
conv3-512	conv3-512	conv3-512	co
	conv1-512	conv3-512	co
			co
maxpool			
conv3-512	conv3-512	conv3-512	co
conv3-512	conv3-512	conv3-512	co
	conv1-512	conv3-512	co
			co
maxpool			
FC-4096			
FC-4096			
FC-1000			
soft-max			

[Karpathy, CS231n]

INPUT: [224x224x3] memory: $224 \times 224 \times 3 = 150K$ params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: $224 \times 224 \times 64 = 3.2M$ params: $(3 \times 3 \times 3) \times 64 = 1,728$

CONV3-64: [224x224x64] memory: $224 \times 224 \times 64 = 3.2M$ params: $(3 \times 3 \times 64) \times 64 = 36,864$

POOL2: [112x112x64] memory: $112 \times 112 \times 64 = 800K$ params: 0

CONV3-128: [112x112x128] memory: $112 \times 112 \times 128 = 1.6M$ params: $(3 \times 3 \times 64) \times 128 = 73,728$

CONV3-128: [112x112x128] memory: $112 \times 112 \times 128 = 1.6M$ params: $(3 \times 3 \times 128) \times 128 = 147,456$

POOL2: [56x56x128] memory: $56 \times 56 \times 128 = 400K$ params: 0

CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800K$ params: $(3 \times 3 \times 128) \times 256 = 294,912$

CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800K$ params: $(3 \times 3 \times 256) \times 256 = 589,824$

CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800K$ params: $(3 \times 3 \times 256) \times 256 = 589,824$

POOL2: [28x28x256] memory: $28 \times 28 \times 256 = 200K$ params: 0

CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400K$ params: $(3 \times 3 \times 256) \times 512 = 1,179,648$

CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400K$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$

CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400K$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$

POOL2: [14x14x512] memory: $14 \times 14 \times 512 = 100K$ params: 0

CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100K$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100K$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100K$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$

POOL2: [7x7x512] memory: $7 \times 7 \times 512 = 25K$ params: 0

FC: [1x1x4096] memory: 4096 params: $7 \times 7 \times 512 \times 4096 = 102,760,448$

FC: [1x1x4096] memory: 4096 params: $4096 \times 4096 = 16,777,216$

FC: [1x1x1000] memory: 1000 params: $4096 \times 1000 = 4,096,000$

TOTAL memory: $24M \times 4 \text{ bytes} \sim 93MB / \text{image}$ (only forward! ~ 2 for bwd)

TOTAL params: 138M parameters

ConvNet Configuration			
B	C	D	
13 weight layers	16 weight layers	16 weight layers	19
put (224 × 224 RGB image)			
conv3-64	conv3-64	conv3-64	conv3-64
conv3-64	conv3-64	conv3-64	conv3-64
maxpool			
conv3-128	conv3-128	conv3-128	conv3-128
conv3-128	conv3-128	conv3-128	conv3-128
maxpool			
conv3-256	conv3-256	conv3-256	conv3-256
conv3-256	conv3-256	conv3-256	conv3-256
	conv1-256	conv3-256	conv3-256
maxpool			
conv3-512	conv3-512	conv3-512	conv3-512
conv3-512	conv3-512	conv3-512	conv3-512
	conv1-512	conv3-512	conv3-512
maxpool			
conv3-512	conv3-512	conv3-512	conv3-512
conv3-512	conv3-512	conv3-512	conv3-512
	conv1-512	conv3-512	conv3-512
maxpool			
FC-4096			
FC-4096			
FC-1000			
soft-max			

[Karpathy, CS231n]



INPUT: [224x224x3] memory: $224*224*3=150K$ params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: $224*224*64=3.2M$ params: $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] memory: $224*224*64=3.2M$ params: $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] memory: $112*112*64=800K$ params: 0

CONV3-128: [112x112x128] memory: $112*112*128=1.6M$ params: $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] memory: $112*112*128=1.6M$ params: $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] memory: $56*56*128=400K$ params: 0

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] memory: $28*28*256=200K$ params: 0

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] memory: $14*14*512=100K$ params: 0

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] memory: $7*7*512=25K$ params: 0

FC: [1x1x4096] memory: 4096 params: $7*7*512*4096 = 102,760,448$

FC: [1x1x4096] memory: 4096 params: $4096*4096 = 16,777,216$

FC: [1x1x1000] memory: 1000 params: $4096*1000 = 4,096,000$

Note:

Most memory is in
early CONV

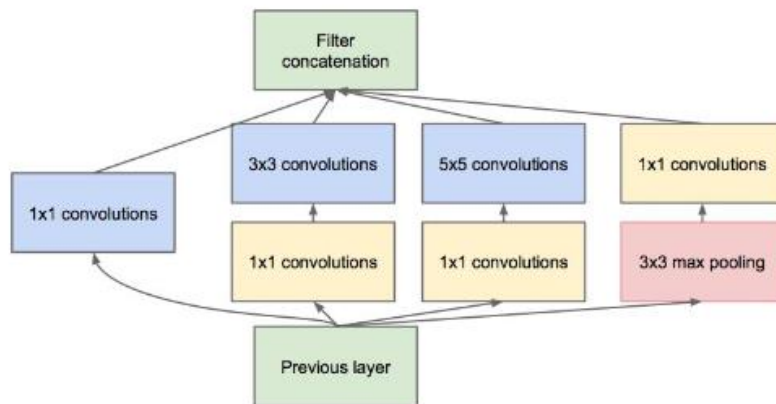

Most params are
in late FC

TOTAL memory: $24M * 4 \text{ bytes} \approx 93MB$ / image (only forward! $\sim *2$ for bwd)

TOTAL params: 138M parameters

Case Study: GoogLeNet

[Szegedy et al., 2014]




Inception module

ILSVRC 2014 winner (6.7% top 5 error)

[Karpathy, CS231n]

Case Study: ResNet *[He et al., 2015]*

ILSVRC 2015 winner (3.6% top 5 error)



2-3 weeks of training
on 8 GPU machine

at runtime: faster
than a VGGNet!
(even though it has
8x more layers)

(slide from Kaiming He's recent presentation)

[Karpathy, CS231n]



Thank you!
Any questions?